Ipcha Mistabra – Institutionalized Contradiction as a Verification Primitive for Epistemically Robust Text Reviews

A Trialectic Questioning Protocol with Reproducible Artifacts and IP-Aware Publication

Oliver Baer
nyxCore Systems Research Division
research@nyxcore.cloud

March 2026

Executive Summary

Safety-critical decisions frequently depend on textual artifacts—specifications, policies, architecture documents, and threat models. Yet text-based reviews, whether manual or LLM-assisted, remain susceptible to confirmation bias, hallucination, and insufficient self-validation. This paper introduces Ipcha Mistabra, a verification protocol that institutionalizes contradiction as a mandatory review phase: a *Proponent* extracts claims and produces an initial assessment; an *Ipcha Agent* (contradictor) systematically constructs counter-hypotheses and demands evidence; an *Auditor* consolidates findings, enforces evidence binding, and renders release-gate decisions. We define formal objects (Claim, Evidence, Finding), the *Ipcha Score* as a divergence-based measure of epistemic correction, provide implementable pseudocode with subfunctions, and outline a reproducible experiment design. The manuscript integrates artifact and IP constraints for Zenodo/DOI publication and double-blind submissions.

Abstract

Background: Reviews of security and compliance text artifacts are central to established frameworks (e.g., BSI IT-Grundschutz, NIST SSDF), yet remain heterogeneous and error-prone in practice [1, 2].
Problem: Modern text-based assistants can over-confirm user assumptions (sycophancy) and, without independent counter-review, contribute to a false sense of security [3]. Existing approaches—including formal inspections, Constitutional AI, and multi-agent debate—address partial aspects but do not combine mandatory adversarial roles, auditable artifacts, and evidence-bound gate rules into a single implementable protocol (see Section 2).
Objective: We introduce a methodology that operationalizes contradiction as a mandatory verification stage, enabling systematic falsification or evidence-based support of claims and assumptions in text.
Method: Ipcha Mistabra defines a trialectic workflow comprising *Proponent* (thesis), *Ipcha Agent* (antithesis), and *Auditor* (synthesis). The protocol extracts atomic claims, generates adversarial challenge questions and counter-hypotheses, requires evidence, and produces auditable findings and remediation plans.
Evaluation (plan): We outline a within-subject experiment comparing a single-pass baseline with Ipcha Mistabra across multiple artifact classes (placeholder datasets), with metrics (recall/precision, risk-weighted residual errors, effort), inter-rater reliability, and a statistical analysis plan (paired tests, mixed models, FDR control).
Expected results & limitations: Higher finding coverage and lower residual errors at increased effort; limitations include ground-truth cost, correlated model failures when roles share a foundation model, and evidence-dependent validation.

Reproducibility/IP: The paper contains a Zenodo-ready artifact manifest and summarizes IP constraints (DE/EU/US) as well as Zenodo metadata publicity [4, 5, 6].

Keywords: security analysis; hardening; text review; adversarial questioning; reproducibility; research artifacts; IP strategy; Ipcha Score.

ACM CCS Concepts (suggested):

- Security and privacy → Software security engineering.
- Software and its engineering → Software verification and validation.
- Computing methodologies → Natural language processing; Artificial intelligence.

1 Introduction

Security-relevant defects often originate early: in requirements, architecture texts, policies, and threat models. Frameworks such as BSI IT-Grundschutz and NIST SSDF emphasize repeatable, documented security practices including reviews and evidence trails [1, 2]. In practice, however, text reviews are inconsistent: reviewers tend to focus on plausibility rather than falsification, and LLM-assisted reviews can confirm user assumptions even when incorrect (sycophancy) [3].

Most mitigation strategies target training-time or inference-time behavioral alignment. These approaches improve average behavior but leave a deeper architectural problem unaddressed: the same system instance is often responsible for both proposing and implicitly validating its own reasoning.

1.1 Research Gap

Several existing approaches address overlapping aspects of this problem: formal inspection methods enforce structured review [7], Constitutional AI introduces rule-based self-critique [8], and multi-agent debate generates adversarial counter-arguments [9]. However, none of these approaches combine all of the following in a single implementable protocol: (i) extraction of atomic claims from text artifacts, (ii) mandatory adversarial contradiction as a distinct role with auditable outputs, (iii) evidence binding between claims and external authority documents, and (iv) a formal gate rule over findings with severity-weighted blocking. Table 1 provides a systematic feature comparison.

1.2 Contributions

1. **Protocol:** Ipcha Mistabra as a trialectic verification primitive for text artifacts.
2. **Formalization:** Definitions (claim, assumption, evidence, finding), pseudocode with specified subfunctions.
3. **Ipcha Score:** A divergence-based quantitative signal for epistemic correction, with specified embedding model and known limitations (Section 4).
4. **Evaluation (plan):** Measurement concept and statistical analysis plan (placeholders for data/results).
5. **Reproducibility & IP:** Artifact manifest, Zenodo-ready metadata, and IP constraints [4, 5, 10, 6].

2 Related Work

2.1 Formal Inspection and Structured Review

Fagan inspection [7] introduced role-separated code and document review (author, moderator, inspector, reader) with defined phases (overview, preparation, meeting, rework, follow-up). The method demonstrated significant defect-detection improvement but is manual, synchronous, and does not formalize evidence binding or produce machine-auditable artifacts. Ipcha Mistabra inherits the principle of role separation but extends it with formal claim/evidence structures and automated gate rules.

2.2 Document-Based Security Processes

BSI IT-Grundschutz describes a systematic approach to information security in which documentation and security review play central roles [1]. NIST SP 800-218 (SSDF) specifies high-level practices for reducing vulnerability risk across the SDLC [2]. These standards motivate reviews but define process guidance rather than an executable verification protocol.

2.3 Failure Modes of LLM-Assisted Text Analysis

Sycophancy denotes the tendency of models to confirm user opinions even at the expense of truth [3]. Perez et al. showed that model-written evaluations can systematically discover such behavioral failures [11]. For safety-relevant text reviews this is especially critical, because "confirmation" can be mistakenly interpreted as "validation."

2.4 Constitutional AI and Rule-Based Self-Critique

Constitutional AI (CAI) [8] introduces a set of principles against which a model critiques its own outputs. This is a single-model approach: the same instance generates and evaluates, which does not structurally prevent correlated blind spots. CAI also lacks auditable, persistent finding artifacts and has no formal gate mechanism. SPARROW [12] extends this with human-authored rules and targeted red-teaming, but remains a training-time intervention rather than a runtime verification protocol.

2.5 Multi-Agent Debate

Du et al. demonstrated that multi-agent debate can improve factuality and reasoning by forcing explicit counter-arguments [9]. However, debate-based setups typically lack: (i) an explicit audit role with synthesis authority, (ii) evidence binding against external sources, and (iii) a formal gate that blocks output based on finding severity. Ipcha Mistabra operationalizes adversarial pressure as a verifiable protocol (roles, artifacts, gate rules) rather than a prompt-based heuristic.

2.6 Adversarial Testing and Red-Teaming

Perez et al. [11] and the MITRE ATLAS framework provide taxonomies and tooling for adversarial testing of AI systems. These focus on model-level vulnerability discovery rather than on text-artifact review. Ipcha Mistabra addresses a different surface: it reviews the *outputs* (documents, specifications) rather than testing the *model itself*.

2.7 Reproducibility and Artifact Policies

Security venues increasingly require and formally evaluate artifacts. USENIX Security 2026 requires open artifact sharing *at submission time* (with justified exceptions) and ties acceptance to continued availability [13, 14]. ACM CCS 2026 requires that artifacts be accessible during

Feature	Fagan	CAI	Debate	ATLAS	**IM Protocol**
Role separation	✓	–	partial	–	✓
Mandatory contradiction	partial	partial	✓	–	✓
Evidence binding	–	–	–	–	✓
Auditable finding artifacts	–	–	–	partial	✓
Formal gate rule	–	–	–	–	✓
Runtime (not training-time)	✓	–	✓	✓	✓
Authority grounding	–	partial	–	✓	✓
Domain-agnostic design	✓	✓	✓	–	✓*

Table 1: Feature comparison of related approaches. *Evaluated in the security domain; generality is a design goal, not yet empirically validated across domains (see Section 6.2).

double-blind review (anonymous URLs/credentials) [15]. ACM additionally defines standardized artifact badges [16].

2.8 Systematic Positioning

3 Method

3.1 Problem Formulation

Given a text artifact T (e.g., specification, policy) and context K (domain, threat model, norms/authority documents). The objective is to produce R, an auditable report containing findings F, evidence bindings, and a remediation plan; optionally with a gate decision (release/block).

3.2 Formal Definitions

Definition 1 (Claim). A *claim* c is an atomic, verifiable statement from T (e.g., "data is encrypted in transit"). The set of all claims is $C = \{c_1, \ldots, c_n\}$.

Definition 2 (Assumption). An *assumption* a is a prerequisite for c that is not explicitly evidenced in T (e.g., "a valid certificate chain is in place"). $A = \{a_1, \ldots, a_m\}$.

Definition 3 (Evidence). *Evidence* e is a traceable attestation (text passage, standard section, test log, code/configuration proof) supporting or refuting a claim. $E = \{e_1, \ldots, e_k\}$.

Definition 4 (Finding). A finding is a tuple

$$f = (id, c, sev, rationale, E_f, remediation, status)$$

with evidence subset $E_f \subseteq E$, severity sev, and status $\in \{\text{accepted}, \text{refuted}, \text{needs-evidence}\}$.

Definition 5 (Gate Rule). A gate rule $\mathcal{G}(F) \in \{0, 1\}$ decides whether release/approval is permitted, e.g.:

$$\mathcal{G}(F) = 1 \iff \nexists\, f \in F : \ (f.status \neq \text{refuted}) \ \wedge \ (f.sev \geq \text{high}).$$

3.3 What Makes This a Protocol, Not a Prompt Chain

The term "protocol" is used deliberately. Three structural properties distinguish Ipcha Mistabra from prompt-based heuristics such as "ask the model to critique itself":

1. **Mandatory execution:** Contradiction is not an optional step that a user may skip; the workflow engine enforces it as a required stage before gate evaluation.

2. **Persistent, auditable artifacts:** Each phase (Proponent, Ipcha, Auditor) produces a typed output (O_p, O_i, O_f) that is stored, versioned, and available for external audit.

Figure 1: High-level architecture of the IPCHA MISTABRA protocol: reasoning layer (Proponent → Ipcha Agent → Auditor) with a verification and memory layer providing authority grounding, epistemic memory, and workflow persistence.

3. **Formal gate mechanism:** The gate rule $\mathcal{G}$ operates over structured finding objects, not over free-text opinions; a BLOCK decision prevents release regardless of subjective confidence.

A prompt chain may implement similar reasoning steps, but without enforcement, persistence, and gating it remains advisory rather than binding.

3.4 Protocol Overview (Trialectic)

IPCHA MISTABRA employs three roles (see Figure 1):

- **Proponent (P):** extracts claims/assumptions, produces initial review O_p and initial findings F_p.
- **Ipcha Agent (I):** generates strongest counter-narratives, counter-hypotheses, and evidence demands; produces O_i and F_i.
- **Auditor (A):** consolidates, deduplicates, prioritizes; enforces evidence; renders gate decision; produces final O_f.

3.5 Step-by-Step Protocol

Step 0: Scope Fixation. Define objective, artifact class, threat model, and authority documents (e.g., BSI/NIST references) [1, 2].

Step 1: Claim Extraction. Segment T; extract C and A; normalize to atomic statements.

Step 2: Proponent Pass. Generate O_p and findings F_p with evidence requirements.

Step 3: Ipcha Pass (Contradiction). Falsify: counterexamples, contradictions, missing prerequisites, norm/policy conflicts; generate O_i, F_i.

Step 4: Audit Resolution. Merge $F = \text{merge}(F_p, F_i)$; status decisions; remediation plan.

Step 5: Gate. Apply $\mathcal{G}$ (release/block).

Step 6: Metrics. Compute review metrics (e.g., finding counts, effort; Ipcha Score as divergence indicator).

3.6 Top-Level Pseudocode

Algorithm 1 IPCHA MISTABRA: Institutionalized Contradiction for Text Artifact Review

Require: Text artifact T, context K, gate rule $\mathcal{G}$
Ensure: Report R, findings F, gate decision g

```
(C, A) ← ExtractClaims(T)                                  ▷ Algorithm 2
(O_p, F_p) ← ProponentReview(T, K, C, A)
(O_i, F_i) ← IpchaContradict(T, K, C, A, O_p, F_p)         ▷ Algorithm 3
F ← MergeAndNormalize(F_p, F_i)
(O_f, F) ← AuditAndResolve(T, K, O_p, O_i, F)
g ← G(F)
R ← AssembleReport(T, K, O_f, F, g)
return R, F, g
```

3.7 Subfunction Specifications

Algorithm 2 EXTRACTCLAIMS: Atomic Claim and Assumption Extraction

Require: Text artifact T
Ensure: Claims C, Assumptions A

```
S ← SentenceSegment(T)                                     ▷ Split T into sentence-level units
for each sentence s ∈ S do
    cands ← IdentifyVerifiable(s)                          ▷ Filter for testable propositions
    for each candidate v ∈ cands do
        c ← NormalizeToAtomic(v)                           ▷ One proposition per claim
        C ← C ∪ {c}
        A_c ← InferPrereqs(c, T)                           ▷ What must hold for c?
        A ← A ∪ A_c
    end for
end for
return C, A
```

When instantiated with an LLM, each subfunction maps to a structured prompt with typed output schema. Appendix C provides example prompt templates.

4 Ipcha Score

We define the **Ipcha Score** as a divergence-based measure of epistemic correction between the Proponent output and the final Auditor output:

Algorithm 3 IpchaContradict: Structured Falsification Pass

Require: T, K, C, A, O_p, F_p
Ensure: Contradictor output O_i, contradictor findings F_i

```
F_i ← ∅
for each claim c ∈ C do
    H⁻ ← CounterHypotheses(c, K)          ▷ Strongest counter-narratives
    V ← AuthorityCheck(c, K)              ▷ Norm/policy violations
    M ← MissingEvidence(c, O_p)           ▷ Gaps in Proponent reasoning
    if H⁻ ≠ ∅ or V ≠ ∅ or M ≠ ∅ then
        f_i ← BuildFinding(c, H⁻, V, M)
        F_i ← F_i ∪ {f_i}
    end if
end for
for each assumption a ∈ A do
    ChallengeAssumption(a, K)             ▷ Demand evidence or flag
end for
O_i ← Summarize(F_i)
return O_i, F_i
```

$$IS \;=\; 1 - \cos\big(\text{embed}(A_p),\ \text{embed}(A_f)\big) \tag{1}$$

where A_p is the Proponent output, A_f is the final Auditor output, and embed$(\cdot)$ denotes a shared embedding function.

Embedding model specification. For reproducibility, we fix the reference embedding to `text-embedding-3-large` (OpenAI, 3072 dimensions) as the default instantiation. Any alternative (e.g., Sentence-T5, BGE-M3, Cohere Embed v3) is permissible provided it is declared and held constant within a study. The score is *not* comparable across different embedding models.

Interpretation (see Figure 2).

- **Low score** (< 0.15): the original answer was already robust, or critique was weak.
- **Mid-range** (0.15–0.45): normal dialectical refinement; the expected operating range for useful adversarial review.
- **High score** (> 0.45): substantial correction pressure; important weaknesses were discovered and resolved.

The threshold values above are preliminary and derived from early pilot observations; they require empirical calibration across artifact types and domains before operational use.

Known limitations of the Ipcha Score. The score measures *semantic surface displacement*, not epistemic correction directly. Two failure modes must be acknowledged:

1. **False positive (high score, low correction):** If the Auditor paraphrases the Proponent output without changing substance, the score inflates. Mitigation: complement with structured finding-level diff (count of status changes in F).
2. **False negative (low score, high correction):** A single critical sentence change (e.g., "system is secure" → "system has unmitigated CVE") may barely shift the embedding centroid. Mitigation: compute a *finding-weighted* variant $IS_w = \sum_{f \in F_\Delta} w(f.sev)$ over findings that changed status between Proponent and Auditor pass.

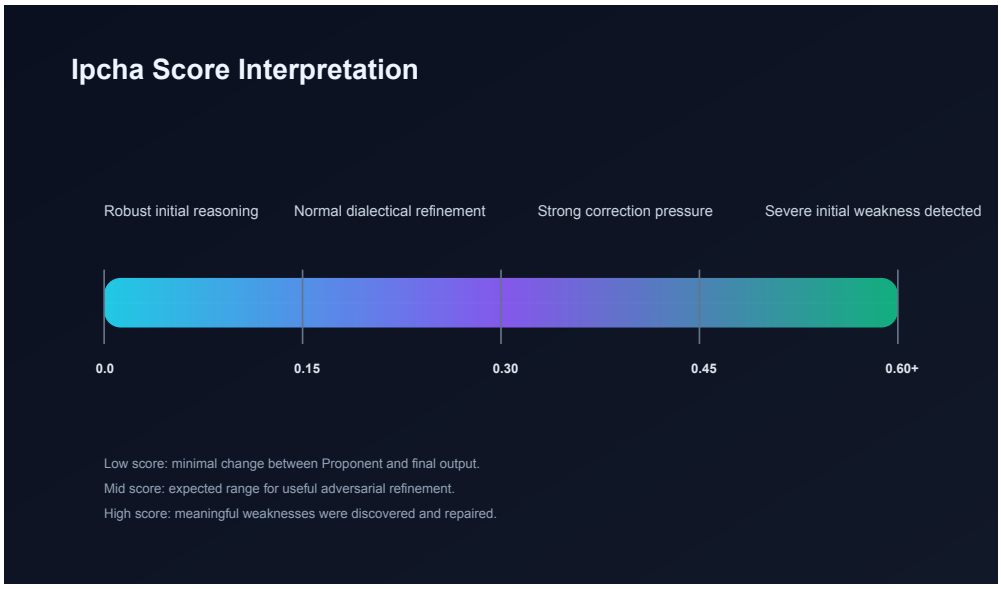

Figure 2: Conceptual interpretation of the Ipcha Score. Thresholds are preliminary; empirical calibration is required. Higher values indicate stronger semantic displacement between Proponent and final validated output.

We recommend reporting *IS* alongside finding-level metrics (recall, precision) and not using it as a sole decision criterion. Future work should validate *IS* against human expert ratings of correction strength (correlation study).

5 Evaluation

5.1 Research Questions

- **RQ1:** Does Ipcha Mistabra increase finding recall/precision relative to single-pass reviews?
- **RQ2:** Does Ipcha Mistabra reduce risk-weighted residual errors and policy violations?
- **RQ3:** What cost/benefit trade-offs arise (time, compute/tokens, human rework)?
- **RQ4:** Does the Ipcha Score correlate with human-judged correction strength?

5.2 Design (Placeholder-Driven)

Design: Within-subject: each artifact is reviewed with baseline (single-pass) and Ipcha Mistabra.
Artifact classes (placeholder): D_{spec}, D_{policy}, D_{threat}, $D_{\text{compliance}}$.
Ground truth: Expert panel or consensus labeling; IRR is reported (see below).

5.3 Metrics

5.4 Hypothesized Direction of Effect

Based on related findings in multi-agent debate [9] and structured inspection [7], we hypothesize the following directional effects. These are *not empirical results*; they define the evaluation target structure for future studies.

Metric	Definition / Measurement
Finding Recall	Proportion of confirmed ground-truth findings discovered
Finding Precision	Proportion of reported findings confirmed by experts
Risk-Weighted Residual Error	Sum of severity for open/high findings after gate
Policy/Compliance Violations	Rate of violations against defined rules/standards
Review Effort	Time (min), compute/tokens (placeholder), rework (min)
Ipcha Score (IS)	Divergence measure per Equation 1
Finding-Weighted Score (IS_w)	Severity-weighted status-change count
IRR	Cohen's κ / Krippendorff's α (depending on setup)

Table 2: Evaluation metrics (numeric values: placeholder).

Metric	Hypothesized direction	Rationale
Error rate	$\downarrow$ (substantial)	Contradiction surfaces hidden flaws
Compliance violations	$\downarrow$ (substantial)	Authority grounding enforces standards
Human review effort	$\downarrow$ (moderate)	Automated pre-screening reduces manual load
Output validity	$\uparrow$	Audit resolution improves final quality

Table 3: Hypothesized directional effects. Magnitude to be determined empirically.

5.5 Statistical Analysis Plan

Paired comparisons: Wilcoxon signed-rank (nonparametric) or paired t-test (if normality is plausible), plus effect size (Cliff's delta / Cohen's d).
Mixed models: Logistic/mixed effects for violations; "artifact" as random effect.
Multiple testing: FDR control (Benjamini–Hochberg).
Power: A priori power analysis (placeholder; depends on expected effect size).
Score validation (RQ4): Spearman rank correlation between IS and expert-rated correction strength.

6 Discussion

6.1 Institutionalized Contradiction as Systems Primitive

The IM Protocol reframes AI reliability as a systems problem rather than a pure model-alignment problem. The approach is *inspired by* fault tolerance in distributed systems, where robustness derives from structured checks between components rather than trust in any single component. We note, however, that the analogy is structural, not formal: the IM Protocol does not provide Byzantine fault tolerance or provable safety/liveness guarantees. Its contribution is to make adversarial challenge a mandatory, logged, and enforceable stage of the workflow—turning dissent into infrastructure rather than leaving it dependent on user prompt quality, developer vigilance, or informal red-teaming.

6.2 Limitations

Correlated model failures. If all three roles (Proponent, Ipcha Agent, Auditor) are instantiated on the same foundation model, they share systematic biases, training-data blind spots, and failure modes. The protocol *reduces* but does not *eliminate* the risk of correlated failures. For example, if a model systematically fails to recognize a class of vulnerabilities (e.g., TOCTOU races), the Ipcha Agent will not construct counter-hypotheses for that class. Mitigations include: (i) model diversification across roles, (ii) authority grounding with external rule sets that explicitly enumerate known vulnerability classes, and (iii) periodic human audit of Ipcha Agent

coverage. The protocol's authority grounding layer (Figure 1) is specifically designed to address this: mandatory documents (legal, compliance, security standards) serve as external oracles that constrain critique beyond the model's intrinsic knowledge.

Domain generality. Ipcha Mistabra is designed with domain-agnostic abstractions (Claim, Evidence, Finding, Gate Rule). However, all examples, metrics, and evaluation targets in this paper are drawn from the security and compliance domain. Applicability to other domains (e.g., medical guideline review, legal contract analysis, scientific peer review) is plausible given the abstract formal structure but remains an untested hypothesis. We explicitly mark domain-agnosticity as a design goal rather than an empirically validated property.

Ground-truth cost. Establishing ground truth for finding recall/precision requires expert panels, which are expensive and domain-specific.

Ipcha Score validity. As discussed in Section 4, the score is a proxy for epistemic correction with known failure modes. It should not be used as a sole evaluation criterion.

6.3 Reproducibility and Artifact Compliance

USENIX Security 2026 requires open artifact sharing at submission; acceptance is tied to continued availability [13, 14]. CCS 2026 requires access to artifacts during double-blind review (anonymous URLs/credentials) [15]. The artifact manifest in Appendix A is designed to satisfy both.

6.4 IP Constraints (DE/EU/Zenodo/Trade Secrets)

Under German law, the state of the art includes everything made publicly available before the relevant date and is novelty-destroying [5]. The EPC defines novelty accordingly (Art. 54) [6]. If parts of the approach are to be treated as trade secrets, GeschGehG requires, among other things, secrecy and appropriate protective measures [10]. Zenodo states that metadata are CC0-licensed and publicly harvestable; metadata are therefore not a suitable repository for secrets [4]. For computer-implemented aspects, note that computer programs "as such" are not patentable; however, computer-implemented inventions may be patentable if they solve a technical problem [17].

IP Note for USENIX 2026. If you wish to preserve DE/EU patent options, "open artifacts at submission time" may conflict with "file before you publish." This tension is one reason the manuscript contains IP sections and redaction strategies [13, 5, 6].

7 Conclusion and Outlook

Ipcha Mistabra institutionalizes contradiction as a verification primitive for safety-relevant text artifacts. We provide a formal specification, pseudocode with specified subfunctions, the Ipcha Score as a quantitative divergence signal with documented limitations, an evaluation design, and an artifact manifest for reproducible research. We argue that institutionalized contradiction should be treated as a foundational primitive for AI workflows where decisions must survive scrutiny beyond surface-level fluency.

Future work includes (i) empirical calibration of the Ipcha Score against human expert ratings (RQ4), (ii) standardized benchmarks across artifact classes, (iii) cross-domain validation (medical, legal, scientific review), (iv) robustness analysis against adversarial text inputs, (v) model-diversification strategies to mitigate correlated failures, and (vi) tooling for semi-automated evidence collection.

Acknowledgments

Funding

No external funding.

Author Contributions

O. Baer: Conceptualization; Methodology; Software; Validation; Writing (Original Draft); Writing (Review & Editing).

Conflict of Interest

The author declares no conflict of interest.

References

[1] BSI-standard 100-2: IT-grundschutz methodology. Federal Office for Information Security (BSI), 2008. DE, PDF. Accessed: 2026-03-13.

[2] NIST SP 800-218: Secure software development framework (SSDF) version 1.1. NIST Computer Security Resource Center, 2022. Accessed: 2026-03-13.

[3] Mrinank Sharma et al. Towards understanding sycophancy in language models. In *International Conference on Learning Representations (ICLR)*, 2024. Placeholder entry; replace with exact bibliographic data if required.

[4] Zenodo policies. Zenodo (CERN), 2026. Metadata licensed under CC0 (except email addresses); metadata is harvestable. Accessed: 2026-03-13.

[5] Patentgesetz (PatG) 3 – state of the art / novelty. Gesetze im Internet (BMJ), 2026. DE. Accessed: 2026-03-13.

[6] EPC article 54 – novelty. European Patent Office, 2020. Accessed: 2026-03-13.

[7] Michael E. Fagan. Design and code inspections to reduce errors in program development. *IBM Systems Journal*, 15(3):182–211, 1976.

[8] Yuntao Bai et al. Constitutional AI: Harmlessness from AI feedback. *arXiv preprint arXiv:2212.08073*, 2022. Placeholder entry.

[9] Yilun Du et al. Improving factuality and reasoning in language models through multiagent debate. In *International Conference on Machine Learning (ICML)*, 2024. Placeholder entry; replace with exact bibliographic data if required.

[10] Trade secrets act (GeschGehG) 2 – definitions. Gesetze im Internet (BMJ), 2026. DE. Definition of trade secret; requires secrecy + reasonable protective measures. Accessed: 2026-03-13.

[11] Ethan Perez et al. Discovering language model behaviors with model-written evaluations. In *Findings of ACL*, 2023. Placeholder entry.

[12] Amelia Glaese et al. Improving alignment of dialogue agents via targeted human judgements. *arXiv preprint arXiv:2209.14375*, 2022. DeepMind SPARROW system.

[13] USENIX security '26: Call for artifacts. USENIX, 2026. Artifact availability tied to (conditional) acceptance; sharing at submission time; exceptions require justification. Accessed: 2026-03-13.

[14] USENIX security '26: Call for papers. USENIX, 2026. Open Science policy; artifacts required at submission time. Accessed: 2026-03-13.

[15] ACM CCS 2026: Call for papers. ACM SIGSAC, 2026. Artifacts at submission time; access instructions for double-blind review. Accessed: 2026-03-13.

[16] ACM artifact review and badging – current. ACM, 2020. Defines badges: Artifacts Evaluated/Available/Results Validated. Accessed: 2026-03-13.

[17] DPMA: Computer-implemented inventions. German Patent and Trade Mark Office, 2026. Computer programs "as such" not patentable; computer-implemented inventions may be patentable if technical problem is solved. Accessed: 2026-03-13.

A Artifact Package (Zenodo-Ready) and File Manifest

File	Format	Purpose
paper_ipcha_mistabra_v1.1.0.pdf	PDF	Main paper (citable)
ipcha_mistabra_zenodo_bundle_v1.1.0.zip	ZIP	Release bundle
SHA256SUMS_v1.1.0.txt	TXT	Checksums (integrity)
README.md	MD	Quickstart, citation note (DOI)
REPRODUCIBILITY.md	MD	Reproduction instructions
MANIFEST.json	JSON	Machine-readable TOC/metadata
LICENSE_DOCS / LICENSE_CODE	TXT	CC BY 4.0 / MIT or Apache-2.0
CITATION.cff	YAML	Citation hints (DOI placeholder)
CHANGELOG.md	MD	SemVer release notes

Table 4: Artifact manifest (suggested file names; contents: placeholder).

B Reproducibility: Commands and Packaging

Dockerfile Skeleton (Placeholder)

```
FROM python:3.12-slim
WORKDIR /app
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt
COPY src/ ./src/
COPY run.sh .
ENTRYPOINT ["bash","./run.sh"]
```

Example CLI

```
unzip ipcha_mistabra_zenodo_bundle_v1.1.0.zip
cd ipcha_mistabra_release_v1.1.0/artifacts/tooling
docker build -t im-tooling:v1.1 .
docker run --rm -v "$(pwd)/../examples:/data" \
  im-tooling:v1.1 \
```

```
  --input /data/example_inputs/spec.txt \
  --context /data/example_inputs/context.json \
  --output /data/example_outputs/report.json
```

C Prompt Templates for LLM Instantiation

The following templates illustrate how the protocol subfunctions can be instantiated with an LLM. They are reference implementations; production deployments should adapt them to the target model and domain.

ExtractClaims Prompt Template

```
SYSTEM: You are a claim extraction agent. Given a text
artifact, extract every atomic, verifiable claim and every
unstated assumption required for that claim to hold.

OUTPUT FORMAT (JSON array):
[
  {
    "id": "C001",
    "claim": "<atomic statement>",
    "source_sentence": "<original text>",
    "assumptions": [
      {"id": "A001", "text": "<unstated prerequisite>"}
    ]
  }
]

Rules:
- One verifiable proposition per claim (atomic).
- Flag implicit prerequisites as assumptions.
- Do NOT include opinions, definitions, or tautologies.
```

IpchaContradict Prompt Template

```
SYSTEM: You are the Ipcha Agent (contradictor). Your role
is structured falsification. For each claim, construct the
strongest counter-hypothesis, check against authority docs,
and identify missing evidence.

INPUT: Claims JSON, Proponent output, Authority documents.

For each claim, produce:
{
  "claim_id": "C001",
  "counter_hypotheses": ["<strongest counter-narrative>"],
  "authority_violations": ["<norm/section violated>"],
  "missing_evidence": ["<what is needed to confirm>"],
  "severity": "critical|high|medium|low",
  "finding_status": "accepted|refuted|needs-evidence"
}
```

```
Rules:
- Assume the claim is FALSE until proven otherwise.
- Ground every objection in authority docs or logic.
- Do NOT generate contrarian noise without substance.
```

D Zenodo Metadata Snippets

Title

```
Ipcha Mistabra: A General Method for Structured Text
Review and Security-Oriented Hardening (v1.1.0)
```

Description

```
Ipcha Mistabra is a verification protocol for structured
text review that extracts claims and assumptions from text
artifacts, adversarially challenges them via a mandatory
contradictor role, and produces evidence-based findings
with formal gate rules.
```

License

```
Documents: CC BY 4.0 | Code/Tooling: MIT or Apache-2.0
```